

Goal of Spring Cloud

- Provide libraries to apply common patterns needed in distributed applications
 - Distributed / Versioned / Centralized Configuration Management
 - Service Registration and Discovery
 - Load Balancing
 - Service-to-service Calls
 - Circuit Breakers
 - Routing
 - ...

Where Does NETFLIX Fit Into All of This?

- Netflix reinvented itself since early 2007
 - Moved from DVD mailing to video-on-demand
 - Once USPS largest first-class customer
 - Now biggest source of North American Internet traffic in evenings.
 - Became Trailblazers in Cloud Computing
 - All Running on Amazon Web Services
 - Chose to publish many general-use technologies as Open-Source projects
 - Proprietary video-streaming technologies are still secret.

Spring and **NETFLIX**

- The Spring Team has always been forward looking
 - Trying to Focus on Applications of Tomorrow
 - Netflix OSS Mature and Battle-Tested; Why Reinvent?
 - Netflix OSS Not Necessarily Easy and Convenient
 - Spring Cloud provides easy interaction
 - Dependencies
 - Annotations

Required Dependencies

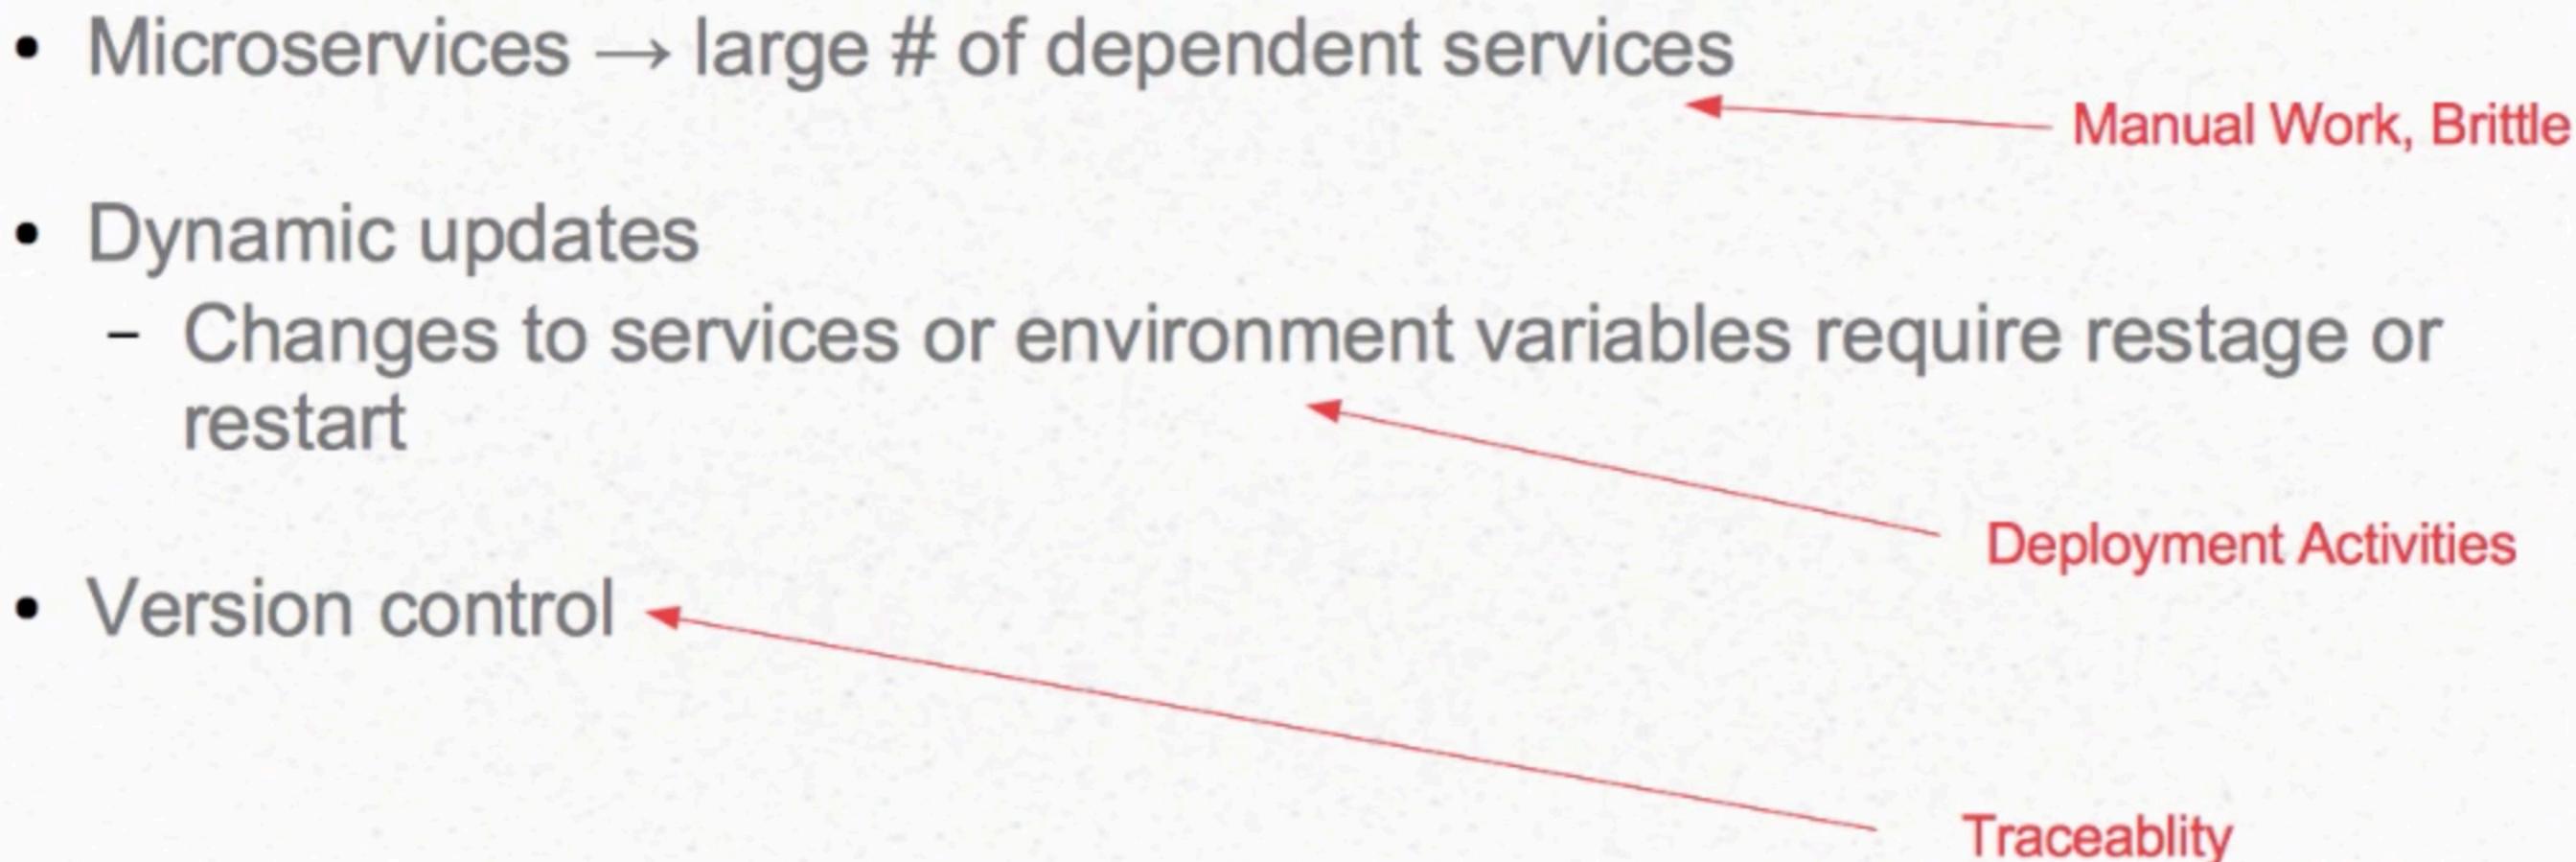
- Replace Spring Boot Parent
 - Spring Cloud projects are based on Spring Boot



```
<parent>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-parent</artifactId>
  <version>Angel.SR4</version>
</parent>

<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-...</artifactId>
</dependency>
```

Other Challenges

- Microservices → large # of dependent services
Manual Work, Brittle
- Dynamic updates
 - Changes to services or environment variables require restage or restartDeployment Activities
- Version controlTraceability

Desired Solution for Configuration

- Platform/Cloud-Independent solution
 - Language-independent too
- Centralized
 - Or a few discrete sources of our choosing
- Dynamic
 - Ability to update settings while an application is running
- Controllable
 - Same SCM choices we use with software
- Passive
 - Services (Applications) should do most of the work themselves by self-registering

Solution:

- Spring Cloud Config
 - Provides centralized, externalized, secured, easy-to-reach source of application configuration
- Spring Cloud Bus
 - Provides simple way to notify clients to config changes
- Spring Cloud Netflix Eureka
 - Service Discovery – Allows applications to register themselves as clients

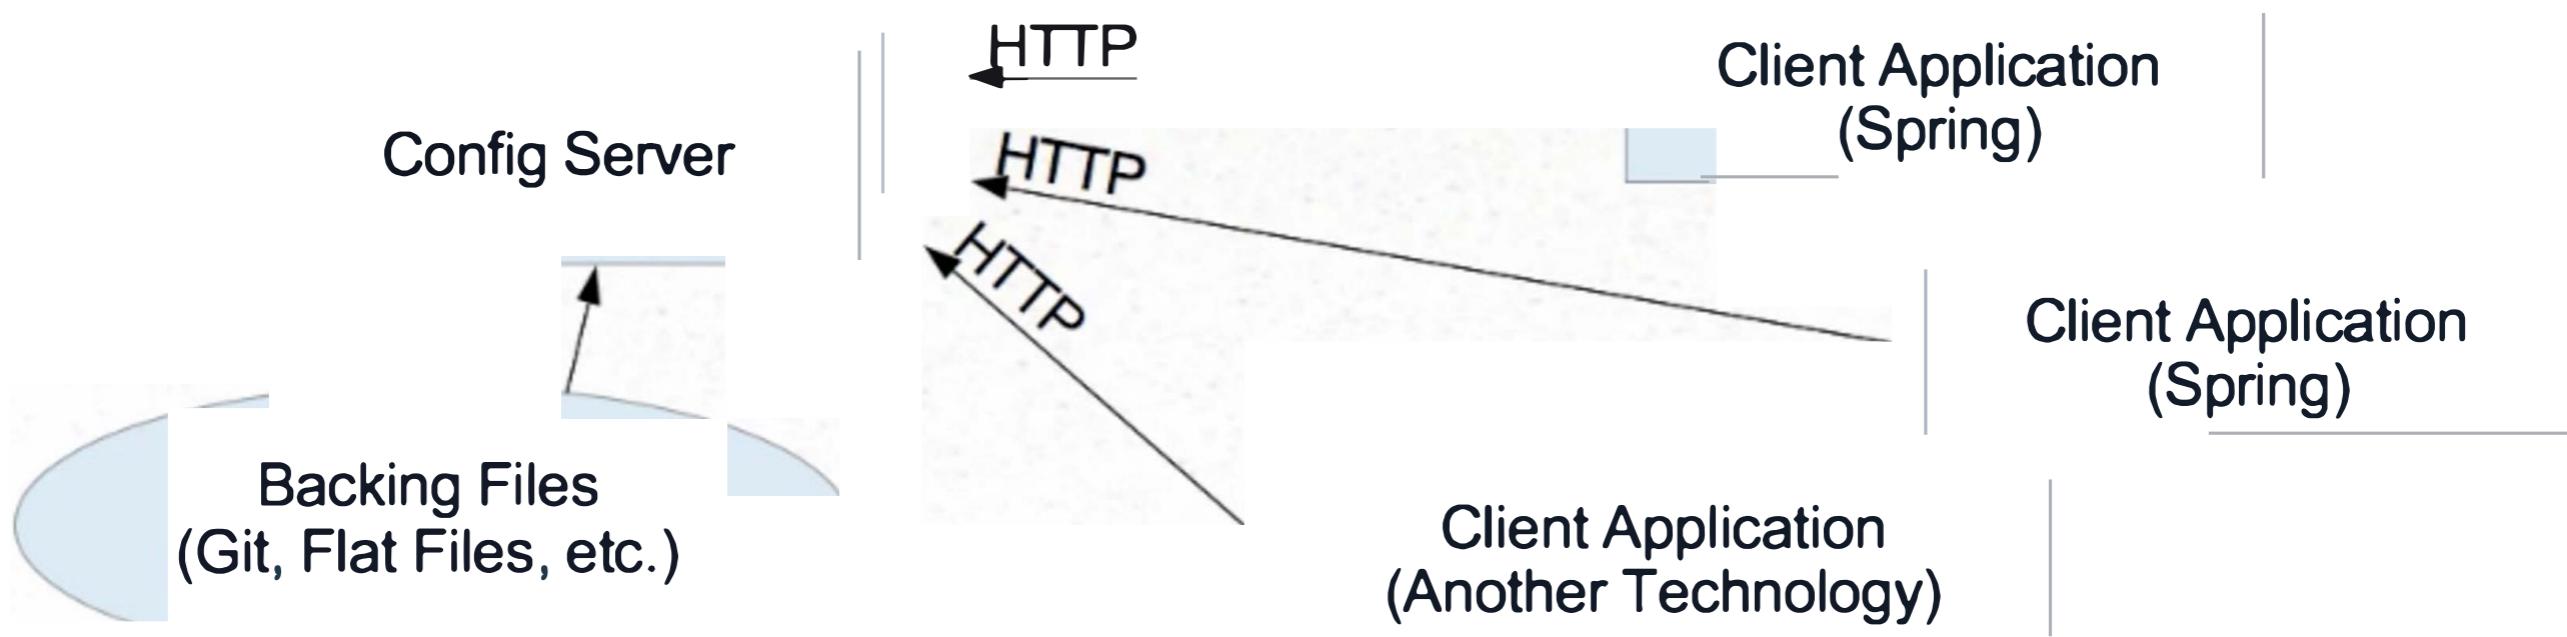
Spring Cloud Config

Designates a centralized server to serve-up configuration information

- Configuration itself can be backed by source control

Clients connect over HTTP and retrieve their configuration settings

- In addition to their own, internal sources of configuration



Spring Cloud Config Server

Include minimal dependencies in your POM (or Gradle)

- Spring Cloud Starter Parent
- Spring Cloud Config Server

```
<parent>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-parent</artifactId>
  <version>Angel.SR4 </version>
</parent>

<dependencies>
  <dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-config-server</artifactId>
  </dependency>
</dependencies>
```

Spring Cloud Config Server –

- application.yml – indicates location of configuration repository

```
spring:  
  cloud:  
    config:  
      server:  
        git:  
          uri: https://github.com/pradeepkrishnamurthy/  
          searchPaths: ConfigData
```

- ...or application.properties

Spring Cloud Config Server –

- Add `@EnableConfigServer`

```
@SpringBootApplication  
@EnableConfigServer  
public class Application {  
  
    public static void main(String[] args) {  
        SpringApplication.run(Application.class, args);  
    }  
  
}
```

EnvironmentRepository - Choices

- Spring Cloud Config Server uses an EnvironmentRepository
 - Two implementations available: Git and Native (local files)
- Implement EnvironmentRepository to use other sources.

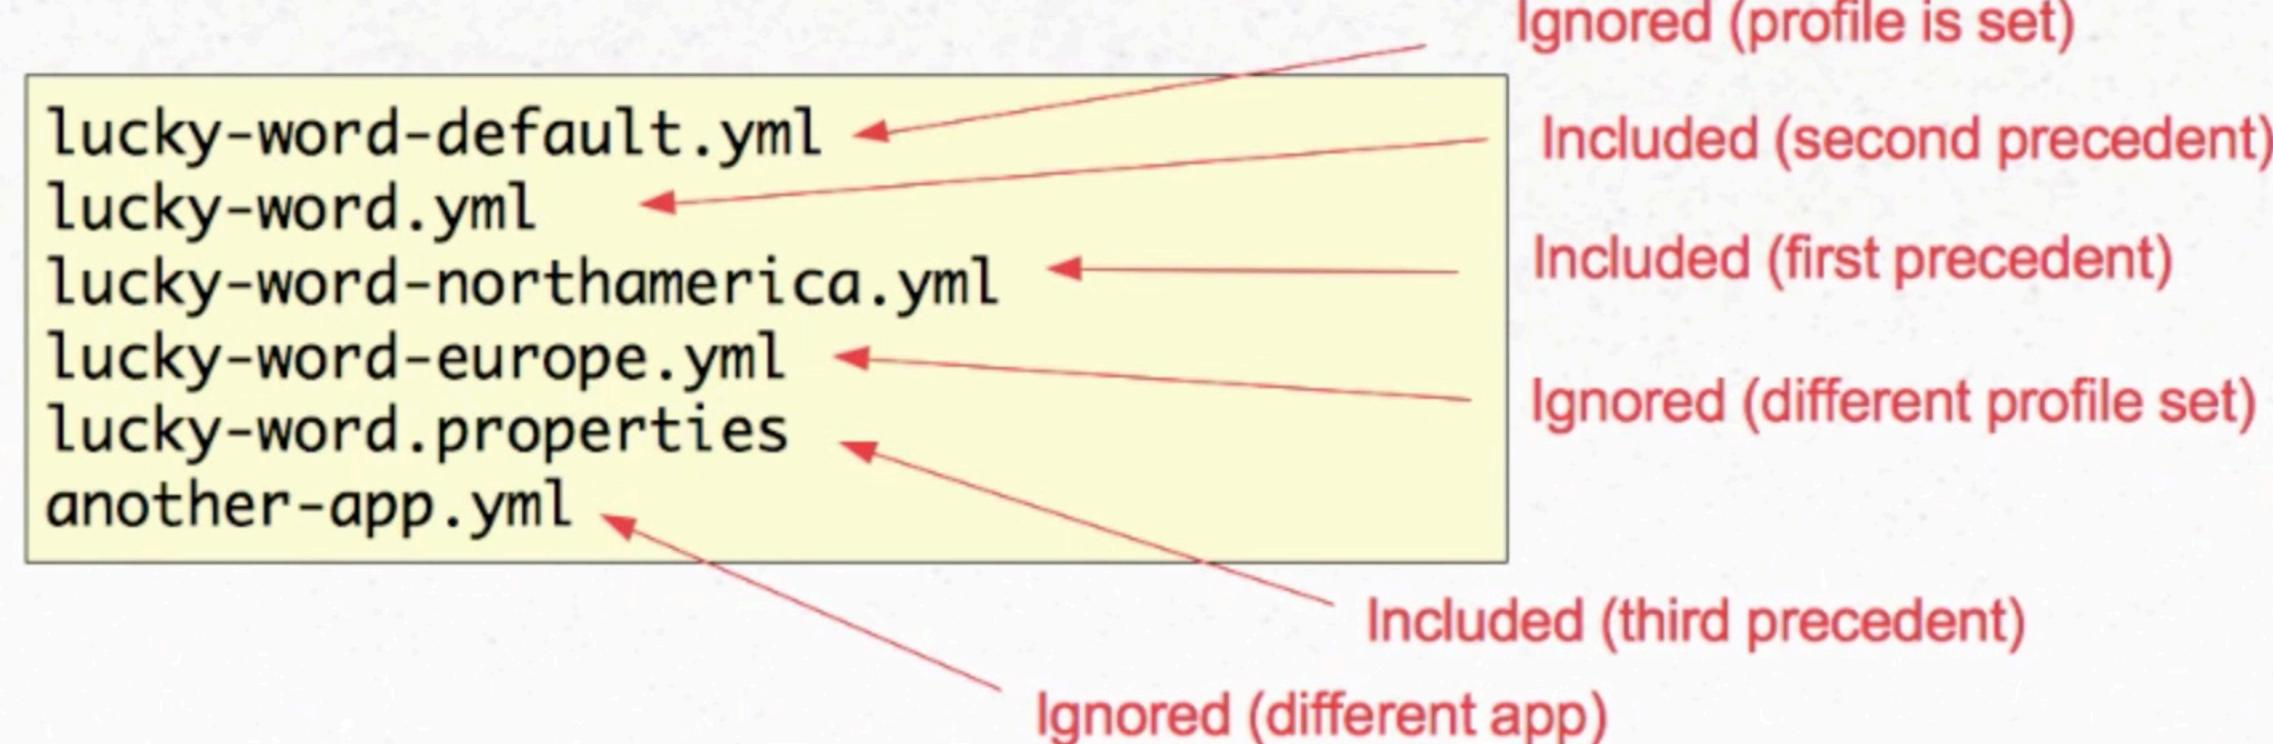
Environment Repository - Organization

- Configuration file naming convention:
 - <spring.application.name>-<profile>.yml
 - Or .properties (yml takes precedence)
 - spring.application.name – set by client application's bootstrap.yml (or .properties)
 - Profile – Client's spring.profiles.active
 - (set various ways)
- Obtain settings from server:
 - `http://<server>:<port>/<spring.application.name>/<profile>`
 - Spring clients do this automatically on startup

Environment Repository – Organization Example

Assume client application named “lucky-word” and profile set to “northamerica”

- Spring client (automatically) requests
 - /lucky-word/northamerica



.yml vs .properties

- Settings can be stored in either YAML or standard Java properties files
 - Both have advantages
 - Config server will favor .yml over .properties

```
# .properties file
spring.config.name=aaa
spring.config.location=bbb
spring.profiles.active=ccc
spring.profiles.include=ddd
```

```
# .yml file
---
spring:
  config:
    name: aaa
    location: bbb
  profiles:
    active: ccc
    include: ddd
```

Profiles

YAML Format can hold multiple profiles in a single file

```
# lucky-word-east.properties  
lucky-word: Clover
```

```
# lucky-word-west.properties  
lucky-word: Rabbit's Foot
```

```
# luckyword.yml  
---  
spring:  
  profiles: east  
  lucky-word: Clover  
  
---  
spring:  
  profiles: west  
  lucky-word: Rabbit's Foot
```

The Client Side

- How Properties work in Spring Applications
 - Spring apps have an Environment object
 - Environment object contains multiple PropertySources
 - Typically populated from environment variables, system properties, JNDI, developer-specified property files, etc.
 - Spring Cloud Config Client library simply adds another PropertySource
 - By connecting to server over HTTP
 - `http://<server>:<port>/<spring.application.name>/<profile>`
 - Result: Properties described by server become part of client application's environment

What if the Config Server is Down?

- Spring Cloud Config Server should typically run on several instances
 - So downtime should be a non-issue
- Client application can control policy of how to handle missing config server
 - `spring.cloud.config.failFast=true`
 - Default is false
- Config Server settings override local settings
 - Strategy: provide local fallback settings.

Service Discovery - Analogy

- When you sign into a chat client, what happens?
 - Client 'registers' itself with the server – server knows you are online.
 - The server provides you with a list of all the other known clients
- In essence, your client has “discovered” the other clients
 - ...and has itself been “discovered” by others

Search people..
Anson Hoy
Eri St. Martin
Hector Virn
Joe Sanantonio
Kevin Croker
Lisa Fernandez
Wei Teh
Adviti a Ban a
Axel Ulrich

Service Discovery

- Microservice architectures result in large numbers of inter-service calls
 - Very challenging to configure
- How can one application easily find all of the other runtime dependencies?
 - Manual configuration – Impractical, brittle
- Service Discovery provides a single 'lookup' service.
 - Clients register themselves, discover other registrants.
 - Solutions: Eureka, Consul, Etcd, Zookeeper, SmartStack, etc.

Eureka – Service Discovery Server and Client

- Part of Spring Cloud Netflix
 - Battle tested by Netflix
- Eureka provides a 'lookup' server.
 - Generally made highly available by running multiple copies
 - Copies replicate state of registered services.
- “Client” Services register with Eureka
 - Provide metadata on host, port, health indicator URL, etc.
- Client Services send heartbeats to Eureka
 - Eureka removes services without heartbeats.

Making a Eureka Server

- Just a regular Spring Boot web application with dependencies and `@EnableEurekaServer`:

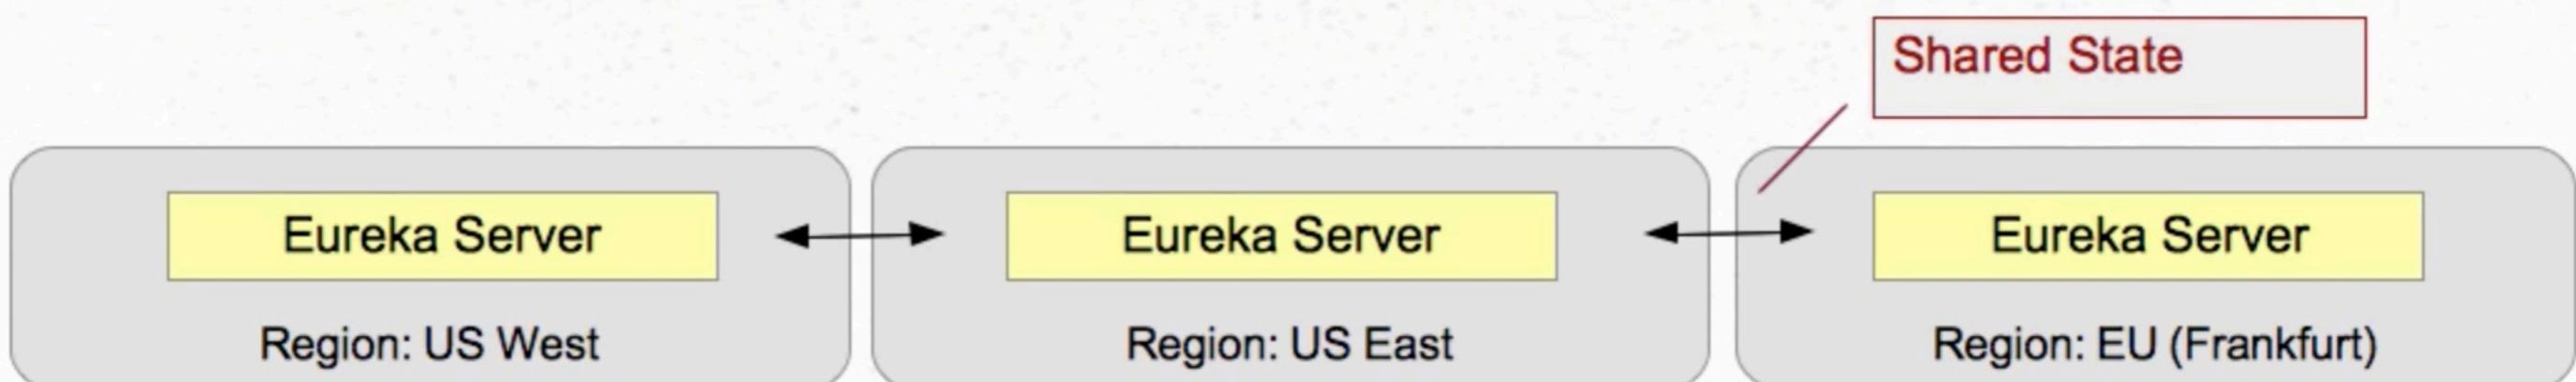
```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-eureka-server</artifactId>
</dependency>
```

```
@SpringBootApplication
@EnableEurekaServer
public class Application {

    public static void main(String[] args) {
        SpringApplication.run(Application.class, args);
    }
}
```

Multiple Servers

- Typically, multiple Eureka servers should be run simultaneously
 - Otherwise, you'll get many warnings in the log
 - Eureka servers communicate with each other to share state
 - Provides High Availability
- Each server should know URL to the others
 - Can be provided by Config Server
 - One server (JAR), multiple profiles



Multiple Servers

Configuration

- Common Configuration Options for Eureka Server:

- See <https://github.com/Netflix/eureka/wiki/Configuring-Eureka> for full list.

Control http port (any boot application)

```
server:  
  port: 8011  
eureka:  
  instance:  
    statusPageUrlPath: ${management.contextPath}/info  
    healthCheckUrlPath: ${management.contextPath}/health  
    hostname: localhost  
  client:  
    registerWithEureka: false  
    fetchRegistry: false  
    serviceUrl:  
      defaultZone: http://server:port/eureka/,http://server:port/eureka/
```

Comma separated list

Registering with Eureka

- From a Spring Boot application

- Add Dependency:

```
<dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-starter-eureka-server</artifactId>
</dependency>
```

- Only the location of Eureka itself requires explicit configuration.

```
@SpringBootApplication
@EnableDiscoveryClient
public class Application {
}
```

```
# application.properties
eureka.client.serviceUrl.defaultZone: http://server:8761/eureka/
```

Default fallback.
Any Eureka instance will do
(we usually want several comma-separated URLs)

@EnableDiscoveryClient

- Automatically registers client with Eureka server
 - Registers the application name, host, and port
 - Using values from the Spring Environment.
 - But can be overridden.
 - Give your application a `spring.application.name`
- Makes this app an “instance” and a “client”
 - It can locate other services

Service ID (Eureka VIP)
Corresponds to
`spring.application.name`

```
@Autowired DiscoveryClient client;
public URI storeServiceUrl() {
    List<ServiceInstance> list = client.getInstances("STORES");
    if (list != null && list.size() > 0 ) {
        return list.get(0).getUri();
    }
    return null;
}
```



What is a Zone?

- Eureka server designed for multi-instance use
 - Standalone mode will actually warn you when it runs without any peers!
- Eureka Server does not persist service registrations
 - Relies on client registrations; always up to date, always in memory
 - Stateful application.
- Typical production usage – many Eureka server instances running in different availability zones / regions
 - Connected to each other as “peers”

Which Comes First? Eureka or Config Server?

Config First Bootstrap (default) – Use Config Server to configure location of Eureka server

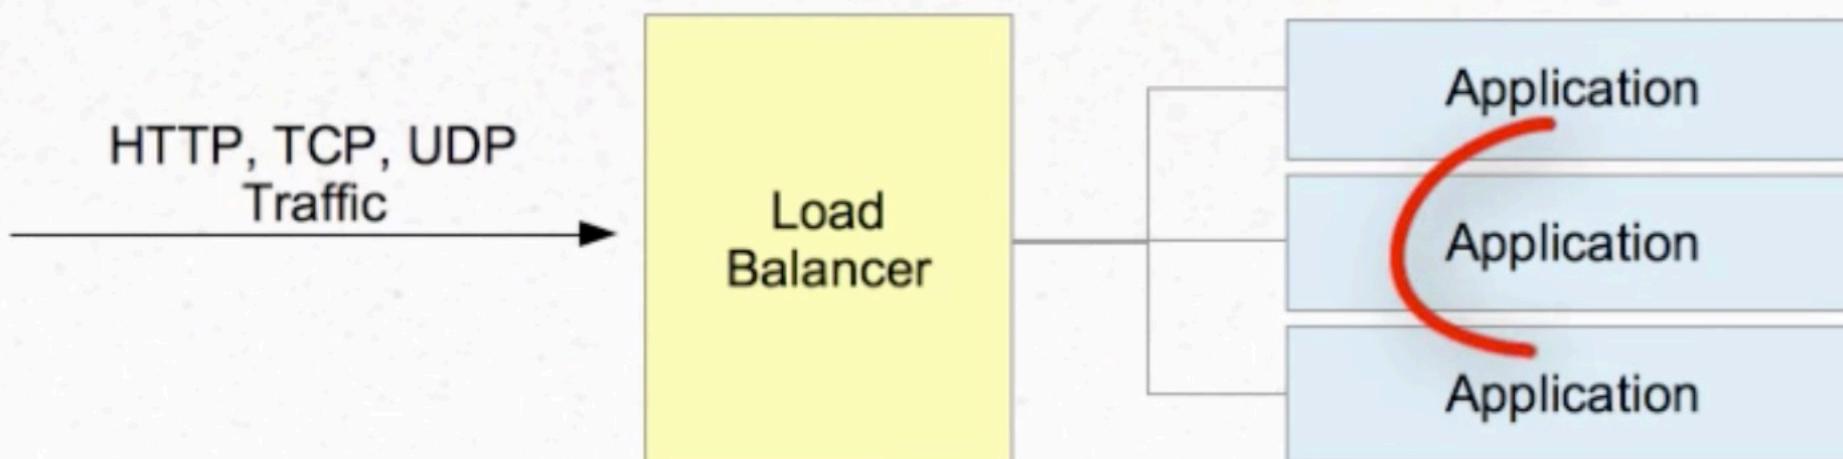
- Implies `spring.cloud.config.uri` configured in each app

Eureka First Bootstrap – Use Eureka to expose location to config server

- Config server is just another client
- Implies `spring.cloud.config.discovery.enabled=true` and `eureka.client.serviceUrl.defaultZone` configured in each app.
- Client makes two network trips to obtain configuration.

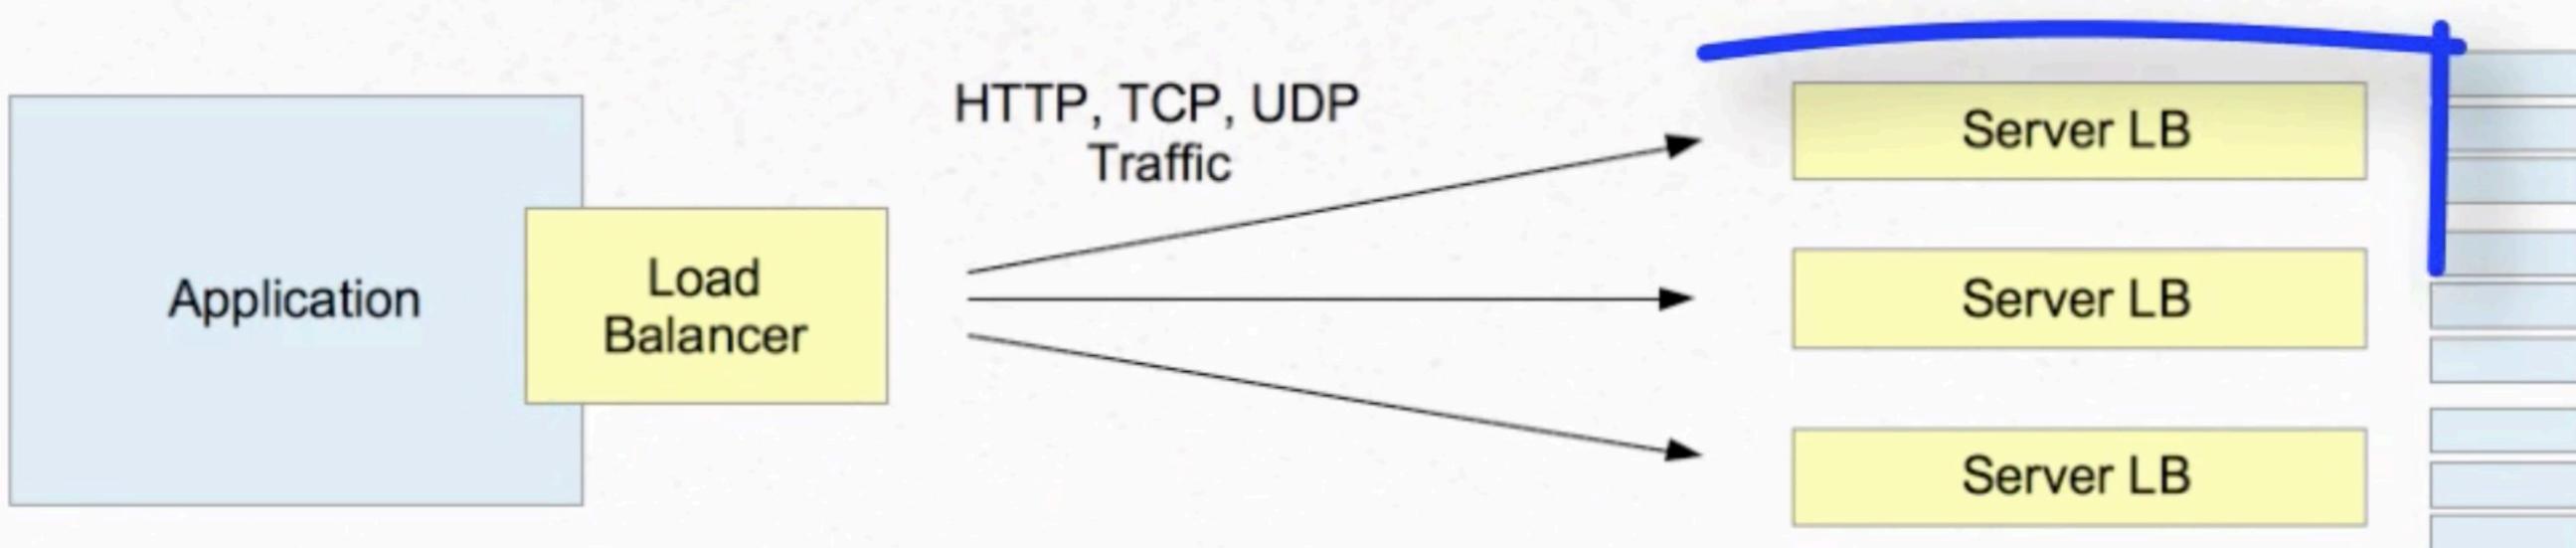
What is a Load Balancer?

- Traditional load balancers are server-side components
 - ▶ • Distribute incoming traffic among several servers
 - Software (Apache, Nginx, HA Proxy) or Hardware (F5, NSX, BigIP)



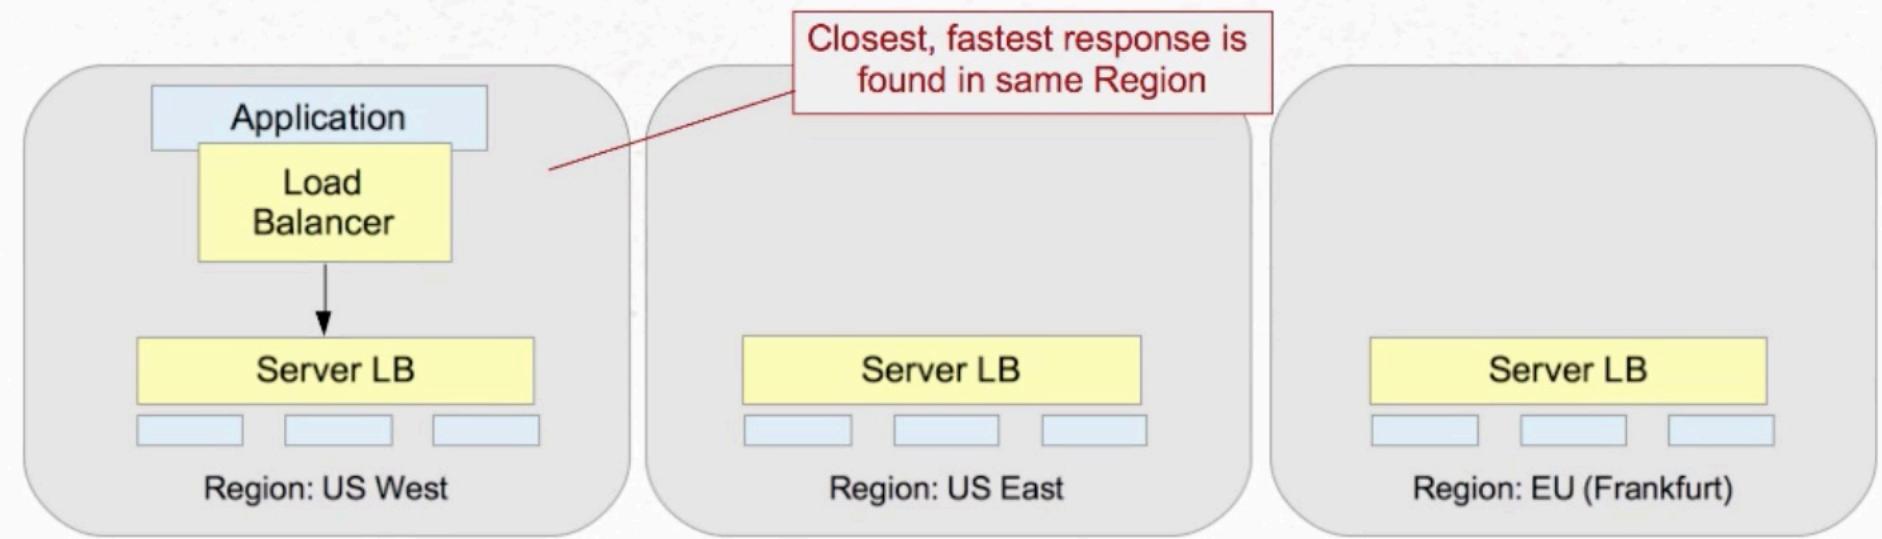
Client-Side Load Balancer

- Client-Side Load Balancer selects which server to call
 - Based on some criteria
 - Part of client software
 - Server can still employ its own load balancer



Why?

- Not all servers are the same
 - Some may be unavailable (faults)
 - Some may be slower than others (performance)
 - Some may be further away than others (regions)



Spring Cloud Netflix Ribbon

- Ribbon – Another part of the Netflix OSS family
 - Client side load balancer
 - Automatically integrates with service discovery (Eureka)
 - Built in failure resiliency (Hystrix)
 - Caching / Batching
 - Multiple protocols (HTTP, TCP, UDP)
- Spring Cloud provides an easy API Wrapper for using Ribbon.

List of Servers

- Determines what the list of possible servers are (for a given service (client))
 - Static – Populated via configuration
 - Dynamic – Populated via Service Discovery (Eureka)
- Spring Cloud default – Use Eureka when present on the classpath.

Example of “Static” server lists

application.yml

“stores” and “products” -
Examples of client-ids

```
stores:  
  ribbon:  
    listOfServers: store1.com,store2.com  
products:  
  ribbon:  
    listOfServers: productServer1.com, productServer2.com
```